

An Efficient Hybrid Peer-to-Peer System for Distributed Data Sharing

Min Yang, *Student Member, IEEE*, and Yuanyuan Yang, *Fellow, IEEE*

Abstract—Peer-to-peer overlay networks are widely used in distributed systems. Based on whether a regular topology is maintained among peers, peer-to-peer networks can be divided into two categories: structured peer-to-peer networks in which peers are connected by a regular topology, and unstructured peer-to-peer networks in which the topology is arbitrary. Structured peer-to-peer networks usually can provide efficient and accurate services but need to spend a lot of effort in maintaining the regular topology. On the other hand, unstructured peer-to-peer networks are extremely resilient to the frequent peer joining and leaving but this is usually achieved at the expense of efficiency. The objective of this work is to design a hybrid peer-to-peer system for distributed data sharing which combines the advantages of both types of peer-to-peer networks and minimizes their disadvantages. The proposed hybrid peer-to-peer system is composed of two parts: the first part is a structured core network which forms the backbone of the hybrid system; the second part is made of multiple unstructured peer-to-peer networks each of which is attached to a node in the core network. The core structured network can narrow down the data lookup within a certain unstructured network accurately, while the unstructured networks provide a low-cost mechanism for peers to join or leave the system freely. A data lookup operation first checks the local unstructured network, and then, the structured network. This two-tier hierarchy can decouple the flexibility of the system from the efficiency of the system. Our simulation results demonstrate that the hybrid peer-to-peer system can utilize both the efficiency of structured peer-to-peer network and the flexibility of the unstructured peer-to-peer network and achieve a good balance between the two types of networks.

Index Terms—Peer-to-peer systems, P2P, structured peer-to-peer, unstructured peer-to-peer, hybrid, overlay networks.

1 INTRODUCTION

LAST few years have witnessed a rapid development of peer-to-peer networks. Research has shown that a large fraction of traffic in the Internet is occupied by peer-to-peer applications [2]. A peer-to-peer (P2P for short) network is a logical overlay network on top of a physical network. Each peer corresponds to a node in the peer-to-peer network and resides in a node (host) in the physical network. All peers are of equal roles. The links between peers are logical links, each of which corresponds to a physical path in the physical network. The physical path is determined by a routing algorithm and composed of one or more physical links. Logical links can be added to the peer-to-peer network arbitrarily as long as a corresponding physical path can be found, that is, the physical network is connected.

The flexibility of the overlay topology and the decentralized control of the peer-to-peer network make it suitable for distributed applications. For example, it can be used for distributed data (file) sharing, where peers announce the data (files) they have and exchange data (files) from each other through a loosely formed peer-to-peer network, or for collaborative Web caching in which Web pages are cached in collaborative peers to reduce network delay for URL requests, or for application layer multicast in which peers

are group members and the peer-to-peer overlay network is a multicast tree. It can also be used for distributed computing which utilizes the idle resources in the network for a huge computing task. Finally, it can be used to provide communication anonymity in which the sender's identity is concealed.

Based on whether a regular topology is maintained among peers, peer-to-peer networks can be divided into two categories: structured peer-to-peer networks in which peers are connected by a regular topology, and unstructured peer-to-peer networks in which the network topology is arbitrary. Structured peer-to-peer networks build a distributed hash table (DHT) on top of the overlay network. The hash table supports efficient data insertion and lookup. Given a key of the data item, the corresponding value of the data item can be inserted or found by transforming the key to a hash value by a hash function. The hash value is the index of the data item and all the hash values form the key space. In DHT, the key space is divided among peers. Each peer is responsible for one partition of the key space. Peers are connected by an overlay network through which the requests of data insertion and lookup are delivered. Structured peer-to-peer networks can provide efficient and accurate query service but need a lot of efforts to maintain the DHT, which makes it vulnerable to frequent peer joining and leaving, also known as *churn*. Churn is a common phenomenon in peer-to-peer networks. Measurement studies of deployed peer-to-peer networks show a high rate of churn [21], [22]. Unstructured peer-to-peer networks organize peers into an arbitrary network topology, and use flooding or random walks to look up data items. Each peer receiving the flooding packets or random walk packets

- The authors are with the Department of Electrical and Computer Engineering, Stony Brook University, Stony Brook, NY 11794. E-mail: {myang, yang}@ece.sunysb.edu.

Manuscript received 14 July 2009; revised 20 Oct. 2009; accepted 29 Oct. 2009; published online 17 Nov. 2009.

Recommended for acceptance by A. Zomaya.

For information on obtaining reprints of this article, please send e-mail to: tc@computer.org, and reference IEEECS Log Number TC-2009-07-0324. Digital Object Identifier no. 10.1109/TC.2009.175.

checks its own database for the data item queried. This approach does not impose any constraint on the network topology. It can perform complex data lookup and support peer heterogeneity. Unstructured peer-to-peer networks are resilient to churn while they usually achieve this goal by sacrificing the data query efficiency and accuracy.

Hence, neither structured peer-to-peer networks nor unstructured peer-to-peer networks can provide efficient, flexible, and robust service alone. The motivation of this paper is to combine the two types of peer-to-peer networks and provide a hybrid solution which can offer efficiency and flexibility at the same time. To achieve this goal, the solution should inherit the advantages of both types in such a way that their disadvantages are minimized.

In this paper, we propose a hybrid peer-to-peer system for distributed data sharing which combines the structured and unstructured peer-to-peer networks. In the proposed hybrid system, a structured ring-based core network forms the backbone of the system and multiple unstructured peer-to-peer networks are attached to the backbone and communicate with each other through the backbone. Data is generated and distributed among the peers. The core structured network provides an accurate way to narrow down the queried data within a certain unstructured network, while the unstructured networks provide a low-cost mechanism for peers to join or leave the system freely.

The main contributions of this paper can be summarized as follows:

- Propose a hybrid peer-to-peer system for distributed data sharing. It utilizes both the efficiency of the structured peer-to-peer network and the flexibility of the unstructured peer-to-peer network, and achieves a good balance between the efficiency and flexibility.
- Give a theoretic analysis on the system performance, and derive quantitative results on the improvement of the join latency and data lookup latency.
- Evaluate the proposed scheme through extensive simulations. Simulation results match the theoretic analysis and show that the proposed scheme achieves the original design goal.

The rest of the paper is organized as follows: In Section 2, we briefly describe some related work. We present the new hybrid peer-to-peer system in Section 3 and analyze its performance in Section 4. We give some enhancements to the hybrid peer-to-peer system in Section 5. In Section 6, we present the simulation results and discuss the performance of the system. Section 7 gives the concluding remark and future work.

2 RELATED WORK

Many peer-to-peer networks have been proposed for different applications in the literature, see, for example, [1], [2], [3], [4], [5], [6], [7], [8], [9], [10], [11], [12], [13], [14], [15]. In this paper, we focus on peer-to-peer networks for efficient distributed data (file) sharing among peers.

The Content Addressable Network (CAN) [10] was proposed to provide a scalable indexing mechanism for file sharing over a large network. As a distributed infrastructure, CAN provides hash-table-like functionality over

Internet-like scales. Both peers and data are hashed to a virtual d -dimensional Cartesian coordinate space. The entire space is partitioned to distinct zones such that each peer is in charge of one zone. Every peer maintains a routing table which holds the IP address of its neighbors in the coordinate space. The data is stored in and retrieved from the peer that owns the zone covering the data. CAN takes advantage of the ordering of the Cartesian coordinate space in the routing algorithm. Packets are forwarded along the straight line connecting the source and the destination in the Cartesian coordinate space. When a new peer joins the system, some existing zone will be split into two zones one of which is assigned to the new peer, and all the related peers need to update their neighbor lists. When a peer leaves the system, a neighboring peer will take over the zone by running a takeover algorithm, and all the related peers need to update their neighbor lists again.

Chord [11] organizes the peers into a circle which is called a chord ring, where each peer is assigned an ID. Peers are inserted into the ring in the order of their IDs. Each peer has two neighbors: successor and predecessor. When a peer joins the system, it first finds the position to insert the new peer. Then, the successor pointers of both the new peer and an existing peer must be changed. The correctness of Chord relies on the fact that each peer is aware of its successor. To guarantee this, each peer maintains a successor list of size r which contains the peer's first r successors. Each data item also has an ID and is stored in a peer such that the ID of the data item is between the ID of the peer and its predecessor. Packets are forwarded along the circle. In order to accelerate the search, each peer maintains a finger table, where each finger points to a peer with a certain distance from the current peer. Chord uses a "stabilization" protocol running in the background to update the successor pointers and finger tables. Compared to CAN, Chord is simpler as the key is hashed in a 1D space.

Gnutella [13] is a decentralized unstructured peer-to-peer network. The network is formed by peers joining the network following some loose rules. There is no constraint on the network topology. To look up a data item, a peer sends a flooding query request to all neighbors within a certain radius. As Gnutella has no requirement on the network topology and data placement, it is extremely resilient to peer joining and leaving the system frequently. However, flooding is not scalable and consumes a lot of network bandwidth. Also, it is difficult to find a rare data item as it has to flood the query request to most of the peers. There has been some work [20] on improving the efficiency when looking up a rare data item.

BitTorrent [14] is a centralized unstructured peer-to-peer network for file sharing. A central server called tracker keeps track of all peers who have the file. Each file has a corresponding torrent file stored in the tracker which contains the information about the file, such as its length, name, and hashing information. When receiving a download request, the tracker sends back a random list of peers which are downloading the same file. When a peer has received the complete file, it should stay in the system for other peers to download at least one copy of the file from it. Since BitTorrent uses a central server to store all the

information about the file and the peers downloading the file, it suffers so called “single point of failure” problem which means that if the central server fails, the entire system is brought to a halt. Note that in some literatures, hybrid peer-to-peer networks were used to refer to the centralized peer-to-peer systems such as BitTorrent. In this paper, “hybrid” is used to refer to the combination of structured and unstructured peer-to-peer network topologies.

YAPPERS [15] combines both structured peer-to-peer networks and unstructured peer-to-peer networks to provide a scalable lookup service over an arbitrary topology. Each peer has an immediate neighborhood (*IN*) which contains all peers within h hops of the peer and an extended neighborhood (*EN*) which contains all peers within $2h + 1$ hops of the peer. Both data keys and peers are hashed to different buckets or colors. Data is stored in the peer in the same color. For a data lookup, a peer first checks the peers in the *IN* in the same color, then these nodes will forward the request to the peers in their *IN*, and so on. Finally, all the peers in the same color will be checked. However, YAPPERS is designed for efficient partial lookup which only returns partial values of data. For a complete lookup, YAPPERS still needs to flood the request to all peers which are in the same color as the data. Compared to YAPPERS, the hybrid peer-to-peer system proposed in this paper can conduct accurate complete lookups in a more efficient way.

In [16], the authors propose a hybrid peer-to-peer system which treats rare and popular data items differently. Some “ultrapeers” form a structured peer-to-peer network which is responsible for caching the rare data. Each ultrapeer has multiple attached leaf peers. Data lookup is first performed through the conventional flooding method. If not successful, the query is reissued to an ultrapeer as a DHT data lookup. It is somewhat similar to the hybrid system proposed in this paper. However, the main difference is that in [16], the structured overlay was used as a supplement for unsuccessful flooded data lookup, while in this paper, the structured overlay is responsible for connecting all the unstructured networks and transmitting queries between them.

3 THE HYBRID PEER-TO-PEER SYSTEM

In this section, we first give an overview of the new hybrid peer-to-peer system we propose, in which a tunable system parameter p_s is defined to characterize the system performance. Then, we describe how to maintain the peer-to-peer network topology when peers join and leave the system. Finally, we describe how to insert and look up data items in the system.

3.1 System Overview

The new hybrid peer-to-peer system is composed of two parts: a core transit network and many stub networks, each of which is attached to a node in the core transit network. The core transit network, called *t-network*, is a structured peer-to-peer network which organizes peers into a ring similar to a chord ring. We call peers in the *t-network* *t-peers*. Each *t-peer* is assigned a peer ID (*p-id*), which is a positive integer. Peers are inserted to the ring in the order of




Fig. 1. Overview of the proposed hybrid peer-to-peer system.

their *p-ids*. Each *t-peer* maintains two pointers which point to its successor and predecessor, respectively. A finger table is also used to accelerate the search. A stub network, called *s-network*, is a Gnutella-style unstructured peer-to-peer network. We call the peers in an *s-network* *s-peers* except for the *t-peer* attached to this *s-network*. The topology of an *s-network* is arbitrarily formed. Each *s-network* is attached to a *t-peer* and this *t-peer* belongs to both the *t-network* and the *s-network*. One thing to mention about the *s-network* is that the topology of an *s-network* is a tree instead of a mesh. We will describe the topology construction in more details later in this section. Fig. 1 shows the overview of the proposed hybrid peer-to-peer system.

We define a system parameter p_s as the proportion of *s-peers* in the system. p_s is a tunable system parameter which has a great impact on the system performance. In particular, when p_s equals to 0, the system degenerates to a ring-based structured peer-to-peer network; when p_s equals to 1, the system becomes a Gnutella-style unstructured peer-to-peer system; and when p_s equals to 0.5, a half of the peers in the system are *t-peers* while the other half of the peers are *s-peers*. By tuning the system parameter, we can observe how the system parameter affects the system performance so that we can select an optimal value for p_s to maximize the performance in different applications. In Section 4, we will give a theoretical analysis on the impact of the system parameter on the system performance for peer joining and data lookup. We will verify our analysis by simulations in Section 6.

The basic idea behind the hybrid peer-to-peer system is that the *t-network* is used to provide efficient and accurate service while the *s-network* is used to provide approximate best-effort service to accommodate flexibility. Peers can join either *t-network* or *s-network* directly. An *s-network* is composed of peers that serve the data of some common properties. A data lookup is confined within an *s-network* if the queried data has the common properties served by the *s-network*. The lookup request is passed around the *s-network* through flooding or random walk. Although flooding may generate a lot of network traffic, it can greatly simplify peer joining and leaving process, which makes the system robust to churn. On the other hand, since the *s-network* contains only a small proportion of the total number of peers, flooding is confined within a small number of peers. When the queried data is served by another *s-network*, the data query request is first forwarded to the *t-network* through the *t-peer* in the *s-network* generating the request. Then, in the *t-network*, the request

will be forwarded along the ring until it reaches the s-network serving the queried data. Finally, in the s-network, the request will be delivered to the s-peers by flooding again. The t-network links all the s-networks together and provides an efficient way to locate the desired s-network. The stableness of the t-network is critical to the system performance because all the communications between different s-networks are through the t-network. As a structured peer-to-peer network, the t-network is vulnerable to churn mainly because the t-network needs to recalculate the pointers in the finger tables whenever a t-peer joins or leaves. However, the hybrid system can effectively reduce the topology maintenance overhead caused by peer joining or leaving. This is because that, on one hand, a large portion of peers join the s-networks directly without disturbing the t-network; and on the other hand, an s-peer can be selected to substitute the leaving t-peer in the same s-network, i.e., the selected s-peer will become a t-peer. In this case, the total number of t-peers is unchanged. Therefore, there is no need to recalculate the pointers in the finger tables, and only a simple update is needed.

In this paper, we focus on applying the hybrid peer-to-peer system to distributed data sharing. A data item is represented by a (*key*, *value*) pair. A *key* is a label or name of the data, such as a file name, while a *value* is the content associated with the key, such as a file. A peer uses operation store (*key*, *value*) to insert the data item into the system and operation lookup (*key*) to obtain the value of the data item. Before performing the store or lookup operations, a peer hashes the data key to an integer *d_id* which is in the same range as *p_id*. As mentioned earlier, s-peers are grouped into different s-networks such that each s-network serves the data of some common properties. In the hybrid system, the *p_ids* of t-peers divide the range of the *d_id* into several segments. Each s-network is responsible for the data whose *d_ids* lie in the same segment. Both store and lookup operations try the local s-network first if the data item is served by the local s-network; otherwise, they turn to the t-network. Thus, such two-tier hierarchy structure can provide efficient lookup in the top tier while maintaining the flexibility in the bottom tier.

Here, we only consider exact-match data lookup. However, it is easy to extend it to support more complex data lookup such as regular-expression-based data lookup. The query message is first flooded within the same s-network. In the meanwhile, it is forwarded to other s-networks through the t-network.

3.2 Peer Join/Leave

Peer-to-peer networks are highly dynamic and autonomic systems in which peers can join or leave the systems at any time. Peers that want to join the system first contact a well-known server to obtain an arbitrary existing peer in the system. The IP address of the server can be acquired by a DNS-like public service. Unstructured peer-to-peer networks process join or leave requests in a more flexible way than structured peer-to-peer networks largely because structured peer-to-peer networks have to maintain the network topology after peers join or leave. Since the hybrid peer-to-peer system contains both structured peer-to-peer network (t-network) and unstructured peer-to-peer network (s-network), the operations of peer joining or leaving the

system are based on which network they belong to. Next, we will discuss these two types of networks separately.

3.2.1 T-Peer Join/Leave

The join operations of t-peers are similar to that in Chord [11]. Each t-peer maintains two pointers that point to its successor and predecessor along the ring, respectively. After the joining peer obtains the arbitrary t-peer, it sends a join request containing its *p_id* to the t-peer. This join request will be forwarded along the ring until it reaches the t-peer such that the *p_id* of the joining peer is between the *p_id* of the t-peer and its successor. The t-peer will initiate a join process and the joining peer is inserted between the t-peer and its successor.

The *p_id* of a new peer is generated at the server. The server has several options to generate the *p_id*. One way is to generate the *p_id* by hashing the IP address of the new peer. Another way is to generate the *p_id* based on the location of the new peer. This can make the peers that are close to each other in the physical network also close to each other in the overlay network. Moreover, the server can generate a random *p_id* for the new peer. However, the *p_id* generation process does not guarantee the uniqueness of *p_id*. In the case of a conflict, the t-peer initiating the join process will generate a new *p_id* which lies in between the *p_id* of itself and its successor. The new *p_id* can be random or simply the midpoint for load balancing purpose.

After the join process completes, the segment of the id space represented by the successor has changed. The peers in the successor's corresponding s-network should transfer part of its data load to the new peer, which is referred to as load transfer. The peers check the data items it stores and transfer all the data items whose *d_id* lie between the *p_id* of the new peer and the predecessor of the new peer.

In a structured peer-to-peer network such as Chord, when a peer leaves the system, all the pointers related to the leaving peer must be updated. These pointers include the successor pointer in its predecessor, the predecessor pointer in its successor, and the finger pointers in all the peers on the Chord ring. The hybrid peer-to-peer system can reduce some update work as it does not need to recalculate new finger pointers. When a t-peer wants to leave the system, it transfers its role to one of the s-peers in its s-network. The selected s-peer will change its role from an s-peer to a t-peer and take over the neighbors and the pointers of the original t-peer. By substituting the t-peer with an s-peer, the number and position of t-peers remain the same. Other t-peers only need to substitute the leaving t-peer with the new t-peer in the finger table.

Sometimes, peers will leave the system without notice due to peer crash. We call it abruptly leaving. To handle abruptly leaving, peers send "HELLO" messages (also called heartbeat messages) to their neighbors periodically. Each peer maintains a timer for each of its neighbors. The time is reset on receiving a "HELLO" message from the corresponding neighbor. Timeout indicates peer crash. The disconnected s-peers will compete to replace the crashed t-peer by sending messages to the server. The server will pick an s-peer to be the new t-peer. The selection can be random or choosing the peer with the smallest IP address.

TABLE 1
Join/Leave Algorithm for t-Peers

```

//Inserting new peer n between peer pre and peer suc
pre.join(n):
  //peer pre checks if ids conflict
  pre.check(n.id); // n.id is the p_id of peer n
  //peer pre sets its successor pointer to peer n
  pre.successor = n;
  //peer n sets its predecessor pointer to peer pre
  n.predecessor = pre;
  //peer n sets its successor pointer to peer suc
  n.successor = suc;
  //peer suc sets its predecessor pointer to peer n
  suc.predecessor = n;
  //peer suc transfers part of its data load to peer n
  suc.loadtransfer(n.id);

n.leave():
  if (s-network != NULL)
    pick a s-peer randomly;
    set s-peer as the new t-peer;
    transfer load from n to the new t-peer;
  else
    //peer pre sets its successor pointer to peer suc
    pre.successor = suc;
    //peer suc sets its predecessor pointer to peer pre
    suc.predecessor = pre;
    //peer n transfers all of its data load to peer suc
    n.loaddump();

pre.check(n.id):
  if (n.id == id) // id is the p_id of peer itself
    //if ids conflict, assign a new id to peer n
    n.id = (id + suc.id) / 2;

suc.loadtransfer(n.id):
  foreach peer in the current s-network
    foreach data in database
      if (data.id <= n.id) // data.id is the d_id of data
        //insert the data item for which peer n is
        //responsible into peer n
        n.insert(data);
        //delete the data item from peer suc
        suc.delete(data);

n.loaddump():
  //move all the data items from peer n to peer suc
  foreach data in database
    suc.insert(data);
    n.delete(data);

```

Table 1 lists the pseudocode for join and leave operations.

3.2.2 S-Peer Join/Leave

Each *s-peer* belongs to an *s-network* and maintains a list of its neighbors. A neighbor can be either an *s-peer* or a *t-peer*. After an *s-peer* acquires the IP address of the random peer, it

adds the peer to its neighbor list. Then, it notifies the random peer to add itself to its neighbor list and the join operation is completed. The *p_id* of the *s-peer* is the same as its neighbor.

It depends on the type of applications that which *s-network* an *s-peer* should be assigned to. In a system with interest-based *s-networks*, as will be discussed in Section 5, the assignment is determined by the interests of the *s-peer*. In other systems, the assignment could be random to distribute the load evenly. The server is responsible for assigning a joining *s-peer* to some *s-network* with a smaller size.

In a Gnutella-style peer-to-peer network, the data lookup is through flooding. The range of flooding is determined by the search radius, that is, the time-to-live (TTL) value of the packet. As the data is distributed around the network randomly, the search radius is critical to the probability of finding the desired data item. For the same topology and the same peer that initiates the search, the longer the search radius, the higher the probability of finding the desired data item, but the longer the latency required. Note that if we add some simple constraints on choosing the random peer when a new peer joins, we may shorten the network diameter, and thus, reduce the search radius without sacrificing the success ratio of finding the desired data item. Next, we will discuss how to add such constraints.

First, we restrict the random peer to be picked to only *t-peers*. Thus, all *s-peers* in the same *s-network* are connected with one *t-peer*. As a result, the diameter of an *s-network* is at most two, and one data lookup can reach all the peers within two hops. The topology of the *s-network* is a star centered at a *t-peer*. Although the data lookup can achieve short latency in such an *s-network*, there is a notable disadvantage that the load is extremely unbalanced. The *t-peer* maintains a long neighbor list while the *s-peer* has a neighbor list with only one neighbor. Each data lookup request has to be forwarded through the *t-peer*. In order to alleviate this problem, we put another restriction on the degree of peers. When the degree of a peer reaches a threshold δ , it passes the join request to one of its neighbors randomly. The join request will be passed until it arrives at a peer whose degree is less than δ . In our simulation, we use this scheme for *s-peer* joining. The new *s-peer* searches from a *t-peer* along a random branch until it finds a peer with a degree less than δ . This peer is called the *connect point* (*cp*) of the new *s-peer*. Besides the neighbor list, each *s-peer* maintains two pointers that store the address of the *t-peer* of the *s-network* and its *cp*. The resulting topology of an *s-network* is a tree. Here, we use a tree instead of a mesh due to bandwidth efficiency consideration. A major drawback of an unstructured peer-to-peer network is that the flooded query messages occupy a lot of network bandwidth. In a mesh network, it is very likely that a peer receives the same query message multiple times from different neighbors. On the other hand, a tree structure guarantees that each peer receives the query message exactly once.

When an *s-peer* leaves the system, it should notify all its neighbors about the leaving. The neighbors then delete it from their neighbor lists. The neighbor whose *cp* is the leaving peer should rejoin the *s-network* by sending a join

request to the t-peer again. The leaving s-peer should also choose a neighbor to transfer the load to.

Again, we need to handle the abruptly leaving when peers crash without notice. To detect and recover from these errors, we still use the periodic “HELLO” messages. Each s-peer periodically broadcasts “HELLO” messages to all its neighbors. “HELLO” messages will reset the timer and timeout indicates peer crash. As the topology of an s-network is a tree with a maximum degree of δ , it is vulnerable to abruptly leaving. When an s-peer crashes, all the peers in the subtree which is rooted at one of the children of the crash peer will be disconnected from the system until the error is detected by the timeout. Shortening the period of the timer can reduce the time the peers are disconnected. However, the drawback is that it generates more network traffic from “HELLO” messages. To overcome this problem, we let each peer respond a data query message with an acknowledgment message. There is a timer for acknowledgment messages. Resetting the timer for acknowledgment messages also resets the “HELLO” timer so that the scheduled “HELLO” message is canceled to save network bandwidth. Now a peer can detect a peer crash by either receiving no “HELLO” message before timeout or receiving no acknowledgment message before timeout. There is another suppress timer. The peers send acknowledgment messages only when the suppress timer is timeout and a data query message is received. This timer is used to suppress the excessive acknowledgment messages when the data query is frequent. In summary, the timers for the acknowledgment messages are used to ensure that the peers respond to link failure or peer crash more quickly when there are data queries. The timer for “HELLO” messages is used to ensure the peers maintain a consistent view of the overlay topology in the absence of data query messages.

Another effect of peer crash is that the load on the peer is lost. As will be seen in the simulation results, we examined this scenario and observed its impact on the data lookup failure ratio.

3.3 Concurrent Join/Leave

Peer-to-peer networks are highly dynamic systems since peers usually are end hosts that are in charge of different individuals or groups. Concurrent joins and leaves are very common and can greatly degrade the performance if not handled carefully.

The concurrency handling in the s-network is simpler than that in the t-network. When an existing peer receives two join requests, it follows the First Come First Serve (FCFS) rule, that is, the second join request will be passed to the next neighbor if the degree of the peer reaches the limit after receiving the first join request. When two or more s-peers leave the system simultaneously, no special action is needed because the disconnected parts will rejoin the s-network, as described in the previous section.

For the t-network, concurrency handling is much more complicated as it involves three t-peers and the topology constraint is strict. The concurrent joins or leaves may lead to an incorrect topology or break the t-network apart if not handled carefully. For example, suppose a t-peer is inserting a new peer, say, x , between itself and its successor. After setting its successor pointer to the new peer, it receives another join request indicating that another new peer, say, y ,




Fig. 2. Concurrent join/leave operation for t-peers.

should be inserted between x and its successor. Thus, the t-peer will pass the join request to its new successor x . However, the join operation of x is not completed, and the successor and predecessor pointers in x are not set correctly. Therefore, the join request of y will not be handled correctly.

We adopt the idea of the concurrency handling in the database system for the concurrent joins or leaves in the t-network. The join and leave requests are sequentialized such that the next request is not processed until the previous request finishes. For a join request, it follows a join triangle, as shown at the left of Fig. 2. When peer pre receives the join request of the new peer, it sets a mutex variable *joining* that indicates some peer is being inserted between peer pre and peer suc . Now peer pre will not accept any leave requests including that from itself. If a new join request comes before the previous request is completed, peer pre will insert the new request to a queue and process the request queue after the previous join request finishes. Then, peer pre sends a packet including the address of peer suc to the new peer. The new peer sets its successor and predecessor pointers to suc and pre , respectively, and sends another packet to peer suc . After receiving the packet, peer suc updates its predecessor pointer and sends a packet back to peer pre . When peer pre receives this packet, it sets its successor pointer to the new peer and continues to process the next join request in the queue. If the queue is empty, it resets the mutex variable.

When a peer leaves the system, it follows the leave triangle which is shown on the right of Fig. 2. When a peer is leaving, it also sets a mutex variable *leaving*. Now the peer will not accept any new join request (if the join request queue is not empty, the peer should process the join request first) and leaving request. Then, the leaving peer sends a packet to peer pre including the address of peer suc . Peer suc sets its successor pointer to peer suc and sends a packet to peer suc including the address of the leaving peer. After receiving the packet, peer suc will check whether the leaving peer included in the packet is what its predecessor pointer is pointing to. Only if they are the same peer, will the peer suc set its predecessor pointer to peer pre and send a packet to the leaving peer to notify the completion of the leaving operation.

3.4 Data Insertion/Lookup

Data is generated and inserted to the system by peers. As mentioned earlier, each s-network is responsible for a range of ID space. The peer generating the data item first hashes the key into this space. If the $d.id$ lies in the range of the current s-network, the data item is inserted to its database and the data insertion is completed. If the $d.id$ does not lie in the range, the data item is sent to the t-peer of the current s-network. Then, it is forwarded

along the t-network until it reaches the t-peer in charge of the ID range covering d_id . Then, the data item is inserted into the database of the t-peer.

Note that although this data placement scheme is simple and easy to implement, the data load between different peers may be imbalanced. Each t-peer corresponds to an s-network. The data generated in all other s-networks except the one a t-peer is in will always be stored in the t-peer; therefore, the data load in t-peers is much heavier than that in s-peers. To alleviate the load imbalance, we spread the data load to the neighbors of the t-peer, i.e., the s-peers connected to the t-peer. When a t-peer receives a data insertion request, it picks a random s-peer from its directly connected s-peers and itself, and then sends a data insertion request to the chosen peer. The random peer will perform the same random data load spreading until the data item is finally inserted to the system. In the performance evaluation section, we will implement both the original data placement scheme and the improved one and study their impacts on the probability density function of the number of data items per peer.

When a peer looks up a specified data item, it first obtains the d_id of the data item by hashing the data key. If the d_id lies in the current s-network, the peer floods lookup packets around the s-network and sets a timer for it. The timer will be reset if the peer receives the data item or expire, which indicates that the data item is not found. The peer may choose to increase the TTL value and the expiration duration of the timer and reflood the lookup packets. If the d_id does not lie in the current s-network, the peer sends a lookup request to the t-peer and also sets a timer for it. Similar to data insertion, the data lookup request is forwarded along the t-network until it arrives at a proper t-peer which will then flood data lookup packets around its s-network. Each peer receiving the lookup request will check its database for data item d_id . If the data item is found in its database, the peer will stop flooding and send the data item to the peer requesting the data item directly.

4 PERFORMANCE ANALYSIS

In this section, we analyze the performance of the proposed hybrid peer-to-peer system and compare it with the structured peer-to-peer networks and the unstructured peer-to-peer networks. In particular, we vary the system parameter p_s , the proportion of the s-peers in the system, and observe how it affects the system performance.

4.1 Performance of the Join Operation

The efficiency of the join operation can be measured by the join latency, which is defined by the time difference between the joining peer sending the join request to the server and the joining peer being inserted into the system.

To give a quantitative analysis on the join latency, we use the number of hops the join request passes to estimate the join latency. Assume that the total number of peers in the system is N . Then, the number of s-peers is $p_s * N$ and the number of t-peers is $(1 - p_s) * N$. For t-peers, the join request is passed along the ring. Thus, the average number of hops passed by the join request is $(1 - p_s) * N/2$. Finger tables can reduce this number to $\log((1 - p_s) * N/2)$. For s-peers, there are two cases. If the selected random peer is

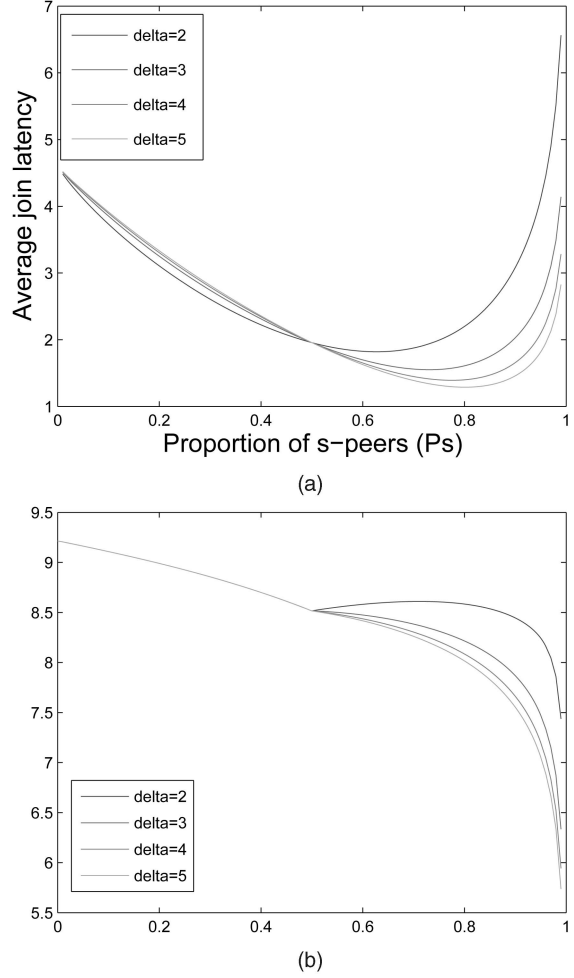


Fig. 3. (a) The average join latency under different δ s. (b) The average lookup latency under different δ s.

always a t-peer, the number of hops passed by the join request is always 1; thus, the join latency is a constant. If we impose the degree constraint when an s-peer joins, the number of hops passed by a join request is determined by the average number of s-peers in an s-network. Recall that an s-network has a tree topology where the t-peer is the root of the tree and the degree of the tree is constrained by δ . The join request is passed from the root to a leaf, which equals the average height of the tree. As mentioned in the previous section, the s-peers are distributed into the s-networks evenly by the server. Therefore, the average number of s-peers in one s-network is $p_s/(1 - p_s)$. Thus, the average number of hops passed by the join request can be calculated by $\log_\delta(p_s/(1 - p_s))$.

We can see that the join latency of t-peers is monotonically decreasing when the system parameter p_s increases, while the join latency of s-peers is monotonically increasing when the system parameter p_s increases (here, we constrain the degree by δ). Therefore, the average join latency of peers is

$$(1 - p_s) * \log((1 - p_s) * N/2) + p_s * \log_\delta(p_s/(1 - p_s)). \quad (1)$$

Fig. 3a plots the curve of the average join latency under different δ values when the system parameter p_s increases

from 0 to 1. We can see that the hybrid peer-to-peer system can achieve a much shorter average join latency than the structure peer-to-peer networks and the unstructured peer-to-peer networks. For example, in the figure, the shortest join latency is achieved when p_s is around 0.7 for $\delta = 2$. We can observe that given system parameter p_s , the larger the degree constraint δ , the shorter the join latency. This is because that a larger δ can reduce the average height of the tree of an s-network, and thus, reduce the number of hops passed by the join request.

4.2 Performance of the Data Lookup Operation

We use the data lookup latency and the data lookup failure ratio to evaluate the data lookup performance. Data lookup latency is defined as the time difference between the time when the peer issues the data lookup request and the time when the peer receives the data. If the peer does not receive the requested data, the timer will expire and this data lookup operation is considered as a data lookup failure. When calculating lookup latency, we only consider the cases of successful lookups. The failed lookups are evaluated by the data lookup failure ratio which is defined by the number of failed data lookups divided by the total number of data lookups.

Before we analyze the performance of the data lookup operation, we define a variable p which is the probability that the requested data item is stored in the same s-network as the s-peer issuing the data lookup. As mentioned earlier, the average number of s-peers in an s-network is $p_s/(1 - p_s)$. Thus, we have $p = p_s/(N * (1 - p_s))$. For a structured peer-to-peer network, each data item is stored at a specific peer. The data lookup request is forwarded to that peer so the lookup failure ratio is 0. For an unstructured peer-to-peer network like Gnutella, the location of a data item is arbitrary, and it uses flooding to perform a best-effort search. The range of the flooding is confined by TTL value (tll) in the flooding packets. To calculate the lookup failure ratio, we need to find the number of peers that are out of the search range of the flooding. In the case that the s-network is constructed without the degree constraint, the lookup failure ratio can be zero if the tll is set to 2. Now we consider the case that the s-network is constructed with the degree constraint δ . In this case, the number of peers that is out of the search range of the flooding is determined by the location of the peer initiating the flooding and the tll . If a t-peer initiates the flooding, the number of peers that are out of the search range is

$$\begin{aligned} & p_s/(1 - p_s) - (1 + \delta + \delta^2 + \dots + \delta^{tll}) \\ &= p_s/(1 - p_s) - (\delta^{tll+1} - 1)/(\delta - 1). \end{aligned}$$

If a leaf s-peer initiates the flooding, the number of peers that are out of the search range is

$$\begin{aligned} & p_s/(1 - p_s) - \sum_{i=2}^{tll/2+1} (\delta^i - 1)/(\delta - 1) \\ &= p_s/(1 - p_s) - (\delta^{2+tll/2} - \delta^2 - (\delta - 1) * tll/2)/(\delta - 1)^2. \end{aligned}$$

We use the midpoint of the two numbers to approximate the average number, which is

$$\begin{aligned} & p_s/(1 - p_s) - (\delta^{tll+2} - \delta^{tll+1} - \delta + 1 \\ &+ \delta^{2+tll/2} - \delta^2 - (\delta - 1) * tll/2)/(2 * (\delta - 1)^2) \\ &\approx p_s/(1 - p_s) - (\delta^{tll+1} * (\delta - 1) + \delta^{2+tll/2} \\ &- (\delta - 1) * tll/2)/(2 * (\delta - 1)^2). \end{aligned} \quad (2)$$

From (2), we can draw the following conclusion: the lookup failure ratio increases if p_s increases while it decreases when tll increases.

We use the number of hops the lookup request packet passes to estimate the data lookup latency. If the s-network is constructed without the degree constraint, the average data lookup latency is

$$p * 2 + (1 - p) * \left(2 + \log\left(\frac{(1 - p_s)N}{2}\right) \right).$$

In the above expression, the first term $p * 2$ represents the average number of hops the flooding packets pass if the data item is stored in the same s-network, and the second term, $(1 - p) * (2 + \log(\frac{(1 - p_s)N}{2}))$ represents the average number of hops the lookup request packet passes if the data item is stored in a different s-network.

Similarly, if the s-network is constructed with the degree constraint δ , the average data lookup latency is

$$\begin{aligned} & p * tll + (1 - p) \\ & * \left(\max\left\{0, \frac{1}{2} \log_\delta\left(\frac{p_s}{1 - p_s}\right)\right\} + tll + \log\left(\frac{(1 - p_s)N}{2}\right) \right), \end{aligned}$$

where $\max\{0, \frac{1}{2} \log_\delta(\frac{p_s}{1 - p_s})\}$ represents the average number of hops the lookup request packet passes from the s-peer starting the lookup to the t-peer, and tll represents the number of hops the flooding packets pass in the s-network which contains the desired data item.

Fig. 3b shows the curve of the average data lookup latency under different δ values when the system parameter p_s increases from 0 to 1. We can see that when p_s is less than 0.5, i.e., the size of an s-network is less than or equal to 1, the data lookup latencies are longer and are the same for different δ . This is because most of the data lookup requests have to pass through the t-network. When p_s is larger, the size of the s-network increases, more data lookup requests can be handled in the local s-network, and the data lookup latency is shortened, though at the expense of a higher lookup failure ratio. Given system parameter p_s , the larger the degree constraint δ , the shorter the data lookup latency. This can also be explained by the fact that a larger δ will have a smaller average height of the tree in an s-network.

To summarize, both the join operation and data lookup operation greatly depend on p_s . The join latency is minimized when p_s ranges from 0.7 to 0.8. This indicates that the hybrid system outperforms the two “pure” systems in terms of join latency. For the data lookup, the selection of p_s is a tradeoff between data lookup latency and lookup failure ratio. In practice, the selection of p_s may have other factors to consider. For example, if the s-networks are interest-based, as discussed in the next section, the data lookup latency largely depends on the TTL instead of p_s as

most data lookup can be conducted within the same s-network. Different applications can adjust p_s to meet their respective requirements.

5 ENHANCEMENTS

In this section, we discuss some enhancements which can make the proposed hybrid peer-to-peer system more flexible and improve the system performance when applied to different applications.

5.1 Supporting Link Heterogeneity

Link heterogeneity is a common phenomenon in networks. In peer-to-peer networks, link heterogeneity refers to the fact that peers have different access link capability. In general, access links include dial-up connections, ADSL, and high-speed cable. The ratio of the link capacity of the fastest link to the slowest link can be more than 1,000.

Due to the imbalance of link capacity, if a peer with high link capacity is receiving messages from a peer with low link capacity, its download speed is upper bounded by the download speed of the low link capacity peer. In other words, the bandwidth of the high link capacity peer is wasted. To maximize the usage of all the link capacities, we connect the peer with higher link capacity in the system to multiple peers with lower link capacities so that the fast peer can download or upload data to the multiple slow peers simultaneously.

As mentioned earlier, the t-peers carry more traffic than the s-peers. Even after we add some restriction to the connect point selection, the load imbalance cannot be removed completely. On the contrary, we can make use of this imbalance for link heterogeneity consideration. Since t-peers are connected to more other peers than s-peers on average, we assign peers with higher link capacities as t-peers while peers with lower link capacities as s-peers. In practice, each peer attaches its link capacity value to the packet sent to the server. Based on the value, the server decides whether the peer is a t-peer or an s-peer.

The idea of link heterogeneity can be applied to the s-network construction as well. When an s-peer joins an s-network, the connect point first checks its *link usage* which is defined as the ratio of the degree to the link capacity of the peer. If the link usage is low enough, the connect point will establish a link between itself and the joining s-peer. If the link usage is high, the connect point will notify the joining s-peer and choose a random branch for it, as discussed in Section 3.2.2.

5.2 Supporting Topology Awareness

Since overlay networks are logical networks on top of physical networks, the overlay links are logical links. Each logic link is composed of one or more physical links. The overlay links are added arbitrarily, as needed. As a result, the topology of the overlay network may be different from the topology of the physical network. Two nodes which are close to each other in the overlay network may be far away in the physical network. Such topology mismatch may increase the *link stress* and degrade the performance, where link stress is defined as the number of copies of a message transmitted over a certain physical link.

One way to solve this mismatch problem is to cluster the peers such that the peers in the same cluster are close to each other with respect to the latency. We can adopt the idea of binning scheme introduced in [17] to construct a topology-aware overlay network. In the scheme, the well-known server is responsible for choosing some peers as landmarks. Each new peer receives the list of landmarks from the server as a response of its join request. The new peer then sends probe messages to the landmarks to learn the distances between them and itself. The landmark peers are listed in an ascending order of distances. The ordered list acts as the coordinate of the new peer in the system. The coordinate is sent to the server, and then the server assigns the new peer a cluster ID based on its coordinate. Peers with the same coordinates form a cluster. Intuitively, peers in the same cluster are close to each other. The server assigns peers in the same cluster to the same s-network. In the case that the number of s-networks is more than the number of clusters, the peers in the same cluster may be assigned to different s-networks. The peers are assigned to several s-networks in a round robin manner such that the numbers of s-peers of different s-networks are balanced.

In addition, every two landmark peers should not be too close to each other. A new peer cannot be a landmark if its coordinate is the same as one of the existing landmark peers. A landmark peer is removed from the landmark list if it has the same coordinate as another landmark peer.

5.3 Interest-Based s-Networks

In previous sections, we assigned the new s-peer randomly to existing s-networks. However, the assignment may be determined by applications. For example, the s-network may be composed of peers that are interested in the same type of data. The data generated in one s-network is looked up mostly by a peer in the same s-network. In this case, a new s-peer can indicate its interest when sending the join request to the server. The server will acknowledge with a t-peer whose s-network matches the interest, and the new s-peer will join that s-network. With the application-specific assignment, s-peers are more likely to generate data stored in their own database and look up data in their own s-network. Hence, the latencies of both data insertion and lookup can be reduced. This way, most of traffic related to data operations is confined within one s-network and the burden of the t-network is alleviated.

Interest-based s-network is also useful for partial/key-word search as well as exact search. The partial search first indicates a field of interest. Then, the partial search is conducted in the corresponding s-network similar to that in other unstructured peer-to-peer networks.

5.4 Bypass Link

As discussed in the previous section, an s-network can be composed of s-peers with the same interest. However, in some applications, an s-peer may be interested in several types of data at the same time. In the hybrid peer-to-peer system, an s-peer can only belong to one s-network at a time. Therefore, we add random bypass links between different s-networks to accelerate the data operations.

Another reason to use bypass links is to alleviate the burden of the t-network. Since the t-network is responsible

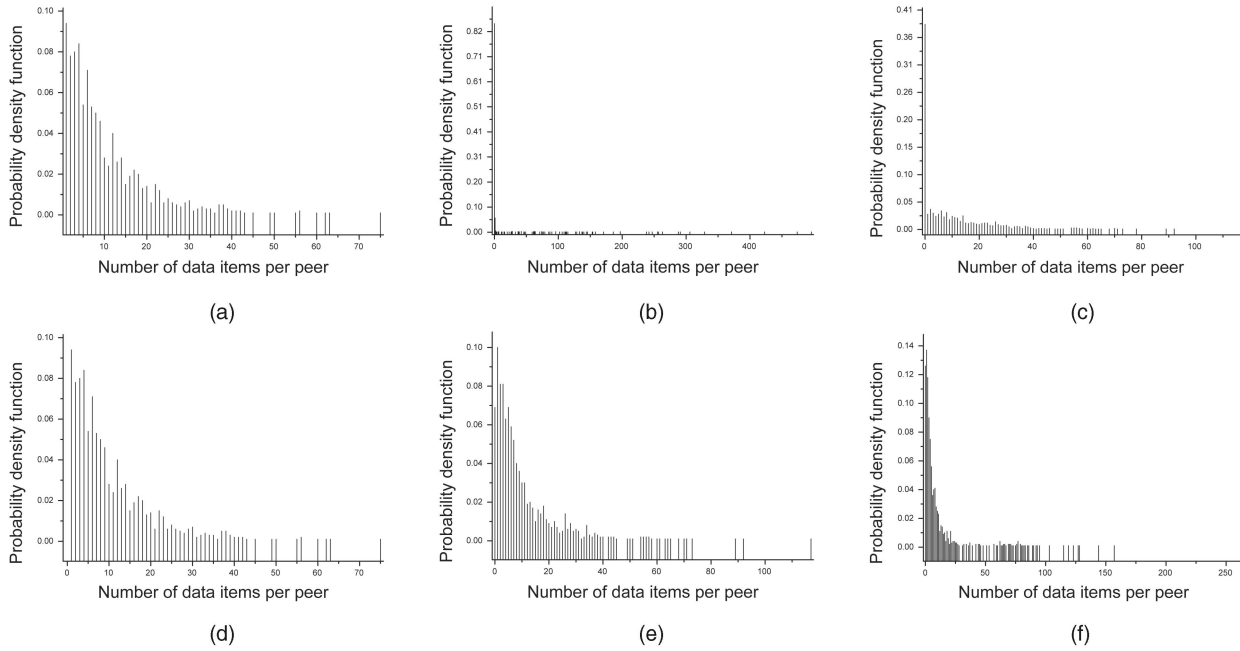


Fig. 4. (a)-(f) Probability density functions of the number of data items per peer under different p_s values for the two data placement schemes.

for transmitting join or leave requests of t-peers and performing all the data operations among s-networks, the t-network links bear much heavier load than that in s-networks. With bypass links, part of the data operations can be diverted from the t-network to the bypass links.

Bypass links can be added according to the following three rules: first, a bypass link cannot be added to a peer unless the degree of the peer is less than the threshold δ ; second, if peer a generates and inserts a data item in peer b , and peer a and peer b are in different s-networks, a bypass link should be added between peer a and peer b ; third, if peer a looks up and finds a data item in peer b , and peer a and peer b are in different s-networks, a bypass link should be added between peer a and peer b .

Peers maintain bypass links in a loose way. Each bypass link is attached with a timer. When the timer expires, the bypass link is deleted. Transmitting a packet through the bypass link will refresh the attached timer. When a peer leaves the system, the bypass links it maintains are lost.

5.5 BitTorrent-Style s-Network

Although the hybrid peer-to-peer system uses the Gnutella network as the prototype of the s-network, it can also easily deploy the BitTorrent-style s-network. In a BitTorrent-style s-network, the t-peer acts as the tracker. When a data item is inserted to the database of an s-peer, it is reported to the t-peer. The t-peer maintains all the information of data items stored in its s-network. A data lookup request is sent to the t-peer directly. T-peer responds to the data lookup with the peer which has the data item. Then, the data item is delivered between the two peers directly. In a BitTorrent-style s-network, the t-peer plays a more important role than that in Gnutella-style s-network and no flooding is needed.

6 SIMULATION RESULTS

In this section, we evaluate the performance of the proposed hybrid peer-to-peer system through simulations. We have

implemented the system on the NS2 simulator. The network topologies used in the simulations are random transit-stub network topologies generated by GT-ITM software [24]. Each network topology is composed of 1,000 nodes, and each node is assigned to be either an s-peer or a t-peer randomly. The ratio between the number of s-peers and the total number of peers is determined by the system parameter p_s . Note that when $p_s = 0$, the hybrid system degenerates to a ring-based structured peer-to-peer network. When $p_s = 1$, the hybrid system becomes a Gnutella-style unstructured peer-to-peer system. In the simulations, we let p_s increase from 0 to 1 so that we can compare the hybrid system with both the structured peer-to-peer network and the unstructured peer-to-peer network. We set the peers with heterogeneous link capacities such that 1/3 of the peers have the highest link capacities, 1/3 of them have the lowest link capacities, and 1/3 of them have the medium link capacities. The highest link capacity is 10 times of the lowest link capacity. The landmarks are predetermined so that they are uniformly distributed around the network. δ is equal to three in the simulations. We will first compare the data distribution under two different data placement schemes discussed in Section 3.4, and then, compare the data lookup failure ratio and data lookup efficiency by tuning the system parameter p_s .

6.1 Data Distribution

As discussed in Section 3.4, when a data item is generated in an s-network and stored in another s-network, it will be forwarded along the t-network until it arrives at a proper t-peer. There are two options for the t-peer: insert the data item to its own database or insert it to the database of a random neighbor. The first option will cause imbalanced data load among peers while the second option can greatly alleviate the imbalance.

Fig. 4 shows the different probability density functions of the number of data items per peer under the two

schemes. In order to display the curves in the center of the figure, the y-axis starts from -0.01 and both x-axis and y-axis are rescaled. Figs. 4a, 4b, and 4c are the probability density functions for the first data placement scheme when p_s is equal to 0, 0.4, and 0.9, respectively. We can see that when $p_s = 0$, more than 50 percent of the peers store less than 10 data items, and the rest of the peers store data items ranging from 10 to 80. When p_s increases to 0.4, the peers without any data occupy 38 percent of the total peers, and the rest of the peers store data items ranging from 0 to 120. When $p_s = 0.9$, the proportion of peers without any data item is 85 percent, and the highest number of data item per peer is more than 500. The peers are evenly distributed along the x-axis. We can observe that with the increase of p_s , the data items are more likely to be stored in t-peers. The data items generated by an s-peer are finally inserted in some t-peer after traveling along the t-network. Most of the s-peers have no data items stored. Figs. 4d, 4e, and 4f show the probability density functions for the second data placement scheme when p_s is equal to 0, 0.4, and 0.9, respectively. The number of peers without any data item drops dramatically. In Fig. 4f, the number of peers with no data item is only around 12 percent compared to 85 percent in Fig. 4e, and 50 percent of the peers store less than 20 data items. We can draw the following conclusion from the figures: when p_s is small, the two schemes can distribute the data items evenly among the peers; when p_s is large, the first scheme stores most of the data items in a few peers while the second scheme can balance the data distribution so that each peer stores a similar amount of data items.

6.2 Lookup Failure Ratio

Lookup failure ratio is defined as the the number of failed data lookups divided by the total number of data lookups. Here, we assume that the data are inserted to the system before it is looked up by some peer.

The lookup failure ratio is affected by the TTL value, the diameter of the network, and the position of the peer initiating the data lookup. Here, we compare the lookup failure ratio of the hybrid peer-to-peer system by tuning both the TTL and p_s .

Fig. 5a shows the lookup failure ratios under different TTL values when p_s increases from 0 to 1. As illustrated in the figure, structured peer-to-peer networks achieve zero lookup failure ratio, while unstructured peer-to-peer networks have a greater than zero lookup failure ratio depending on the TTL value. When $p_s < 0.5$, the lookup failure ratio of the hybrid system is around 0 for all TTL values. Since the number of s-peers is less than the number of t-peers, an s-network has less than one s-peer on the average. Thus, the flooding packets from a t-peer can reach all the s-peers within one hop on the average. When p_s increases, the lookup failure ratio increases exponentially. When TTL = 1 and $p_s = 0.9$, the lookup failure ratio is increased to 18 percent. From the figure, we can also observe that increasing TTL value can reduce the lookup failure ratio dramatically. When $p_s = 0.9$, the lookup failure ratio is four percent if TTL = 4, compared to 14 percent if TTL = 2 and 18 percent if TTL = 1. If we set p_s to less than 0.5, the lookup failure ratio is almost zero which is the same

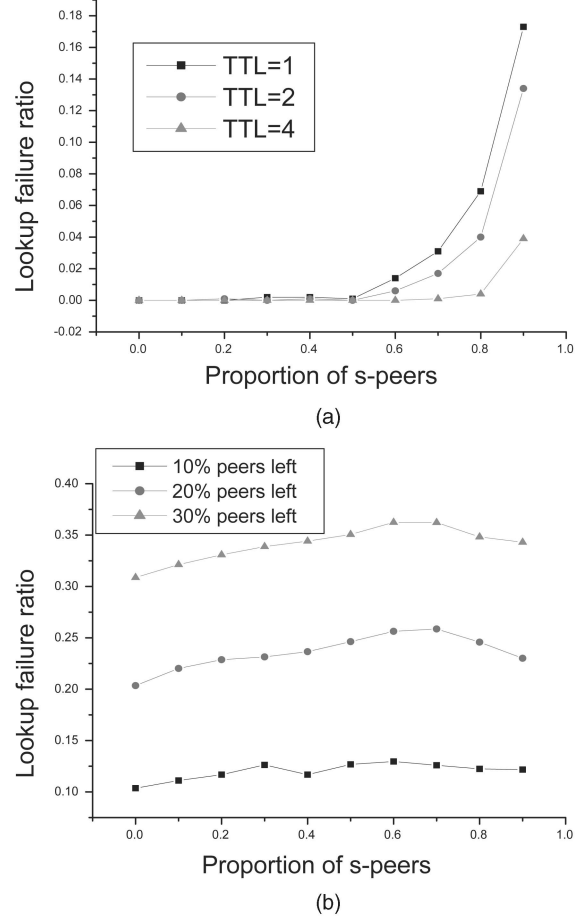


Fig. 5. (a) Lookup failure ratio under different TTL values. (b) Lookup failure ratio when peers crash under different p_s values.

as structured peer-to-peer networks. In the meanwhile, we can minimize the average join latency by setting p_s to around 0.7 (assuming $\delta = 2$), as demonstrated in Section 4. However, the lookup failure ratio will increase accordingly. We can increase the TTL value to lower the lookup failure ratio to some acceptable level. We will discuss this tradeoff later and deduce the optimal p_s value.

We now consider the impact of peer crash on the lookup failure ratio. Peer crash means that the peer has no time to transfer the data load to one of its neighbors. In the simulation, the peers are chosen randomly to leave the system without transferring its data load. Fig. 5b shows the lookup failure ratio when different proportions of peers leave the system. We can see that the lookup failure ratio increases linearly to the number of peers leaving the system. As more peers are leaving the system, the data stored in the peers is no longer available. With the increase of p_s , the lookup failure ratio remains at the same level. This again shows that the improved data placement scheme can distribute the data load among peers evenly regardless of the value of p_s .

Generally speaking, higher p_s value causes higher lookup failure ratio. Increasing TTL value can alleviate the situation greatly. When peers crash, the increase of the lookup failure ratio is proportional to the number of crashed peers and changing p_s has no impact on the lookup failure ratio.

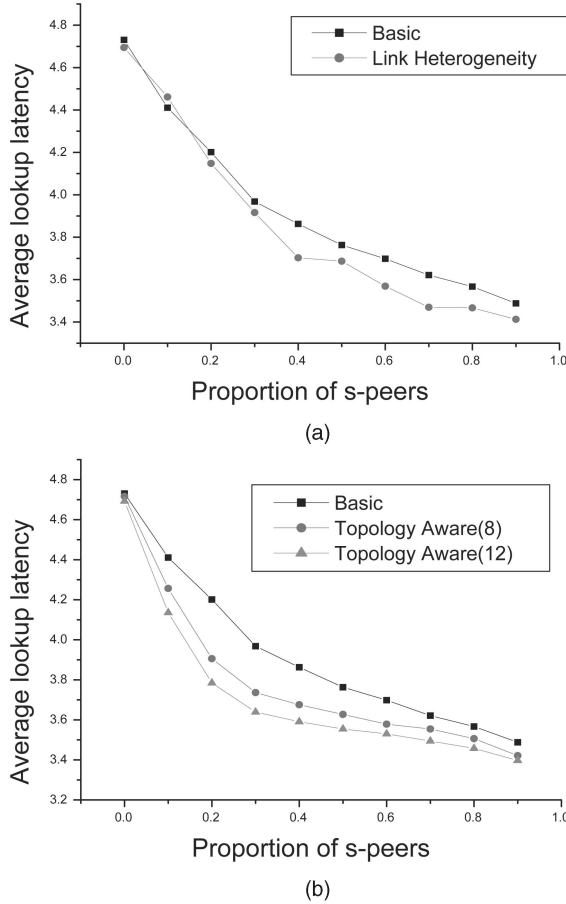


Fig. 6. (a) and (b) Average lookup latency under different p_s values.

6.3 Lookup Efficiency

In an ideal system, a high lookup success ratio should not sacrifice the lookup efficiency. We evaluate the lookup efficiency in terms of time and space. For the time, the lookup efficiency is simply the lookup latency defined earlier. For the space, it is the bandwidth involved in the lookup. We approximate the bandwidth by the number of peers a lookup may contact, denoted as *connum*. The less *connum* is, the less the bandwidth is needed. Here *connum* is the number of peers all the data lookup requests contact during the simulation.

Fig. 6a shows the average lookup latency with and without link heterogeneity consideration under different p_s values. As the figure shows, structured peer-to-peer networks have larger lookup latency than unstructured peer-to-peer networks while a hybrid system has a lookup latency in between. While p_s increases, the average lookup latency decreases more quickly when p_s is small. This is due to the fact that the number of t-peers decreases when p_s increases. A typical data lookup is composed of three steps: an s-peer sends data lookup request to the t-peer; the lookup request is forwarded along the t-network until it reaches a proper t-peer; the lookup request is flooded around the s-network. As both the first and third steps are confined by the TTL value in the lookup packet, the number of hops a lookup experiences is determined by the second step, which is proportional to the total number of t-peers. When p_s is large, the average size of s-networks is large

TABLE 2
Total *connum* Under Different p_s Values

p_s	TTL=1	TTL=2	TTL=4
0	4882354	4882354	4882354
0.1	4384861	4384861	4384861
0.2	3955554	3955556	3955560
0.3	3455618	3455663	3455683
0.4	3014036	3014192	3014231
0.5	2454076	2454480	2454513
0.6	1903119	1903916	1904185
0.7	1393979	1395594	1396464
0.8	919519	924644	929429
0.9	462057	479646	502686

which results in a larger search radius and lookup latency. This latency increase partially offsets the latency decrease so that the slope of the curve is more gradual when p_s is large.

When link heterogeneity is considered, the average lookup latency is further decreased. When p_s is between 0.4 and 0.8, this latency decrease is more visible. The lookup latency decreases by about 20 percent when $p_s = 0.7$.

Fig. 6b shows the average lookup latency with and without topology awareness consideration under different p_s values. We consider two possible numbers of the landmarks: 8 and 12. From the figure, we can see the lookup latencies are the same when $p_s = 0$. When p_s increases, the latency of topology awareness scheme decreases much faster than the basic scheme. The more the landmarks, the less the lookup latency. When p_s approaches 0.9, the three curves tend to merge. This indicates that the topology awareness has little impact when the number of the s-networks is large, because in this case, the peers close to each other are assigned to the same s-network with high probability.

It can be seen that extending the basic scheme to support link heterogeneity or topology awareness can shorten the lookup latency, especially when p_s is around 0.7 for link heterogeneity or around 0.3 for topology awareness. It reveals that the hybrid system is more sensitive to the extensions.

Table 2 shows the total number of peers the data lookups contact under different TTL values. Generally speaking, structured peer-to-peer networks contact more peers to find a data item than unstructured peer-to-peer networks. This does not mean that unstructured peer-to-peer networks are more efficient than structured peer-to-peer networks because the lookup failure ratios of unstructured peer-to-peer networks are much greater than that of structured peer-to-peer networks. *connum* decreases linearly with the increase of p_s . When $p_s = 0.9$, we can see that *connum* is only about 10 percent of that of the structured peer-to-peer networks. This is also due to the three steps of a data lookup. When p_s is small, most of the peers are t-peers. A data lookup will pass a half of the t-network on the average. When p_s is large, the number of t-peers involved in the second step decreases. Although the number of s-peers involved in the first and third steps increases, the increase is very slow as the s-peers are evenly distributed in different s-networks.

We can also observe that the TTL value has no impact on *connum* when p_s is small. However, when p_s is large, the

larger the TTL value is, the larger *connum* is. This is because that a larger TTL value enlarges the range of the flooding packet, thus increasing the number of peers involved. Therefore, we have the following conclusion: the average number of peers a lookup may contact drops linearly as p_s increases. Also, increasing TTL causes slightly greater *connum* only when p_s is larger than 0.5.

From the above discussions, we can see that increasing p_s helps reduce both lookup latency and *connum*, which is consistent with the theoretical analysis in Section 4. It means that the hybrid system can achieve less lookup latency and *connum* than structured peer-to-peer networks. However, a higher p_s value also causes the lookup failure ratio to increase dramatically. Based on the results in Section 4 and the simulations, p_s should be set to an optimal value which is around 0.7 to maximize both the efficiency and the effectiveness of data lookup.

7 CONCLUSIONS AND FUTURE WORK

In this paper, we have proposed a hybrid peer-to-peer system which combines both the structured peer-to-peer network and the unstructured peer-to-peer networks to form a two-tier hierarchy to provide efficient and flexible distributed data sharing service. The top tier is the t-network which is a structured ring-based peer-to-peer network providing efficient and accurate service. The bottom tier is composed of multiple unstructured s-networks which provide approximate best-effort service to accommodate flexibility. By assigning peers to the t-network or the s-network, the hybrid peer-to-peer system can utilize both the efficiency of the structured peer-to-peer network and the flexibility of the unstructured peer-to-peer network and achieve a good balance between them. Our simulation results show that compared to structured peer-to-peer networks, the hybrid system has less lookup latency and *connum*, thus has higher data lookup efficiency. The efficiency can be further increased when considering link heterogeneity and topology awareness. Compared to unstructured peer-to-peer networks, the hybrid system has much lower data lookup failure ratio. By adjusting the system parameter p_s and TTL value, the data lookup failure ratio can be lowered to an acceptable level. Moreover, the join latency can be minimized by selecting appropriate p_s .


Our future work includes how to design a caching scheme for the hybrid peer-to-peer system to improve the system performance. In the case that some extremely popular data are requested by a large amount of peers, the peer hosting the data may be overwhelmed by the large amount of requests. The goal of the caching scheme is to balance the load of the hosting peer when popular data are requested by many peers. The idea is to distribute the load among as many peers as possible so that no peer is overwhelmed. The challenges include how to choose some surrogate peers to redirect the requests to, which data should be cached and how long the data should be cached.

ACKNOWLEDGMENTS

The research was supported in part by the US National Science Foundation under grant numbers CCR-0207999 and CCF-0744234.

REFERENCES

- [1] E.K. Lua, J. Crowcroft, M. Pias, R. Sharma, and S. Lim, "A Survey and Comparison of Peer-to-Peer Overlay Network Schemes," *IEEE Comm. Surveys and Tutorials*, vol. 7, no. 2, pp. 72-93, Mar. 2005.
- [2] S. Sen and J. Wang, "Analyzing Peer-to-Peer Traffic Across Large Networks," *Proc. Internet Measurement Workshop*, Nov. 2002.
- [3] Y. Chu, S. Rao, and H. Zhang, "A Case for End System Multicast," *Proc. ACM Sigmetrics '00*, pp. 1-12, June 2000.
- [4] E. Brosh and Y. Shavitt, "Approximation and Heuristic Algorithms for Minimum Delay Application-Layer Multicast Trees," *Proc. IEEE INFOCOM '04*, Mar. 2004.
- [5] M. Freedman and R. Morris, "Tarzan: A Peer-to-Peer Anonymizing Network Layer," *Proc. Ninth ACM Conf. Computer and Comm. Security*, Nov. 2002.
- [6] S. Banerjee, C. Kommareddy, K. Kar, B. Bhattacharjee, and S. Khuller, "Construction of an Efficient Overlay Multicast Infrastructure for Real-Time Applications," *Proc. IEEE INFOCOM '03*, pp. 1521-1531, Mar. 2003.
- [7] S. Iyer, A. Rowstron, and P. Druschel, "Squirrel: A Decentralized, Peer-to-Peer Web Cache," *Proc. 21st Ann. ACM Symp. Principles of Distributed Computing (PODC)*, 2002.
- [8] A.R. Bharambe, S.G. Rao, V.N. Padmanabhan, S. Seshan, and H. Zhang, "The Impact of Heterogeneous Bandwidth Constraints on DHT-Based Multicast Protocols," *Proc. Ann. Int'l Workshop Peer-to-Peer Systems (IPTPS '05)*, pp. 115-126, Feb. 2005.
- [9] D. Mijolicic, V. Kalogeraki, R. Lukose, K. Nagaraja, J. Pruyne, B. Richard, S. Rollins, and Z. Xu, "Peer-to-Peer Computing," Technical Report HPL-2002-57, HP Labs, Mar. 2002.
- [10] S. Ratnasamy, P. Francis, M. Handley, R. Karp, and S. Shenker, "A Scalable Content Addressable Network," *Proc. ACM SIGCOMM '01*, pp. 161-172, 2001.
- [11] I. Stoica, R. Morris, D. Karger, M.F. Kaashoek, and H. Balakrishnan, "Chord: A Scalable Peer-to-Peer Lookup Protocol for Internet Applications," *IEEE/ACM Trans. Networking*, vol. 11, no. 1, pp. 17-32, Feb. 2003.
- [12] V. Vishnumurthy and P. Francis, "On Heterogeneous Overlay Construction and Random Node Selection in Unstructured P2P Networks," *Proc. IEEE INFOCOM '06*, 2006.
- [13] Gnutella Development Forum, The Gnutella v0.6 Protocol, http://groups.yahoo.com/group/the_gdf/files/, 2009.
- [14] Bittorrent, <http://www.bittorrent.com/>, 2009.
- [15] P. Ganesan, Q. Sun, and H. Garcia-Molina, "YAPPERS: A Peer-to-Peer Lookup Service over Arbitrary Topology," *Proc. IEEE INFOCOM '03*, pp. 1250-1260, 2003.
- [16] B.T. Loo, R. Huebsch, I. Stoica, and J.M. Hellerstein, "The Case for a Hybrid p2p Search Infrastructure," *Proc. Workshop Peer-to-Peer Systems (IPTPS '04)*, pp. 141-150, Feb. 2004.
- [17] S. Ratnasamy, M. Handley, R.M. Karp, and S. Shenker, "Topologically-Aware Overlay Construction and Server Selection," *Proc. IEEE INFOCOM '02*, June 2002.
- [18] H. Zhang, G. Neglia, D. Towsley, and G. Lo Presti, "On Unstructured File Sharing Networks," *Proc. IEEE INFOCOM '07*, pp. 2189-2197, 2007.
- [19] M. Zaharia and S. Keshav, "Gossip-Based Search Selection in Hybrid Peer-to-Peer Networks," *Proc. Ann. Int'l Workshop Peer-to-Peer Systems (IPTPS)*, Feb. 2006.
- [20] Y. Chawathe, S. Ratnasamy, L. Breslau, N. Lanham, and S. Shenker, "Making Gnutella Like p2p Systems Scalable," *Proc. ACM SIGCOMM*, Aug. 2003.
- [21] S. Saroiu, K. Gummadi, and S. Gribble, "A Measurement Study of Peer-to-Peer File Sharing Systems," *Proc. Multimedia Computing and Networking Conf. (MMCN)*, Jan. 2002.
- [22] K.P. Gummadi, R.J. Dunn, S. Saroiu, S.D. Gribble, H.M. Levy, and J. Zahorjan, "Measurement, Modeling and Analysis of a Peer-to-Peer File-Sharing Workload," *Proc. ACM Symp. Operating System Principles (SOSP)*, Dec. 2003.
- [23] M.-T. Sun, C.-T. King, W.-H. Sun, and C.-P. Chang, "Attribute-Based Overlay Network for Non-DHT Structured Peer-to-Peer Lookup," *Proc. Ann. Int'l Conf. Parallel Processing (ICPP)*, Sept. 2007.
- [24] <http://www.cc.gatech.edu/projects/gtitm/>, 2009.



Min Yang received the BEng and MS degrees in computer science from Beijing University of Posts and Telecommunications and Tsinghua University, China, in 1999 and 2002, respectively, and the PhD degree in electrical engineering from Stony Brook University, New York, in 2009. He is currently with FalconStor Software in Melville, New York. His research interests include multicast routing, network coding, and peer-to-peer networks. He is a student member of the IEEE and the IEEE Computer Society.



Yuanyuan Yang received the BEng and MS degrees in computer science and engineering from Tsinghua University, Beijing, China, and the MSE and PhD degrees in computer science from Johns Hopkins University, Baltimore, Maryland. She is a professor of computer engineering and computer science at Stony Brook University, New York, and the director of Communications and Devices Division at New York State Center of Excellence in Wireless and Information Technology. Her research interests include wireless networks, optical networks, high-speed networks, and parallel and distributed computing systems. Her research has been supported by the US National Science Foundation and the US Army Research Office. She has published over 200 papers in major journals and refereed conference proceedings and holds six US patents in these areas. She is currently an associate editor for the *IEEE Transactions on Computers* and a subject area editor for the *Journal of Parallel and Distributed Computing*. She has served as an associate editor for the *IEEE Transactions on Parallel and Distributed Systems*. She has served as a general chair, program chair, or vice chair for several major conferences and a program committee member for numerous conferences. She is a fellow of the IEEE and the IEEE Computer Society. More information about her and her research can be found at <http://www.ece.sunysb.edu/yang>.

► **For more information on this or any other computing topic, please visit our Digital Library at www.computer.org/publications/dlib.**